



Software & Packages


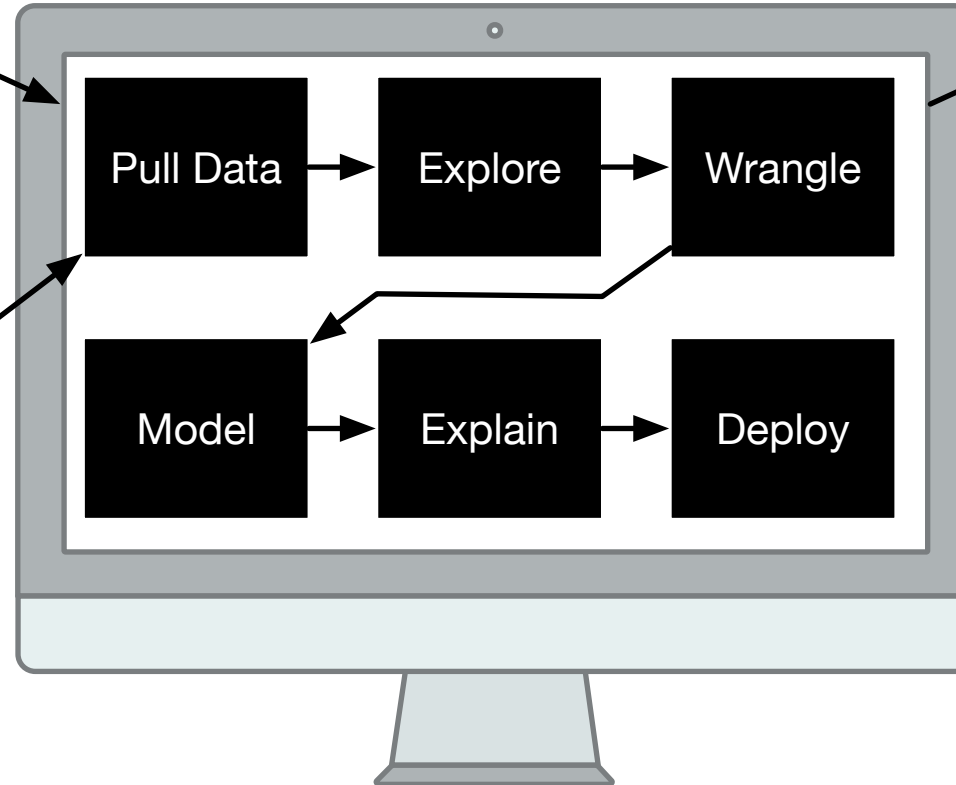
VCS, e.g. Github

Data Scientist Workstation




apt, pip

wget




CI / CD

Download




Data



Production Systems

api



User

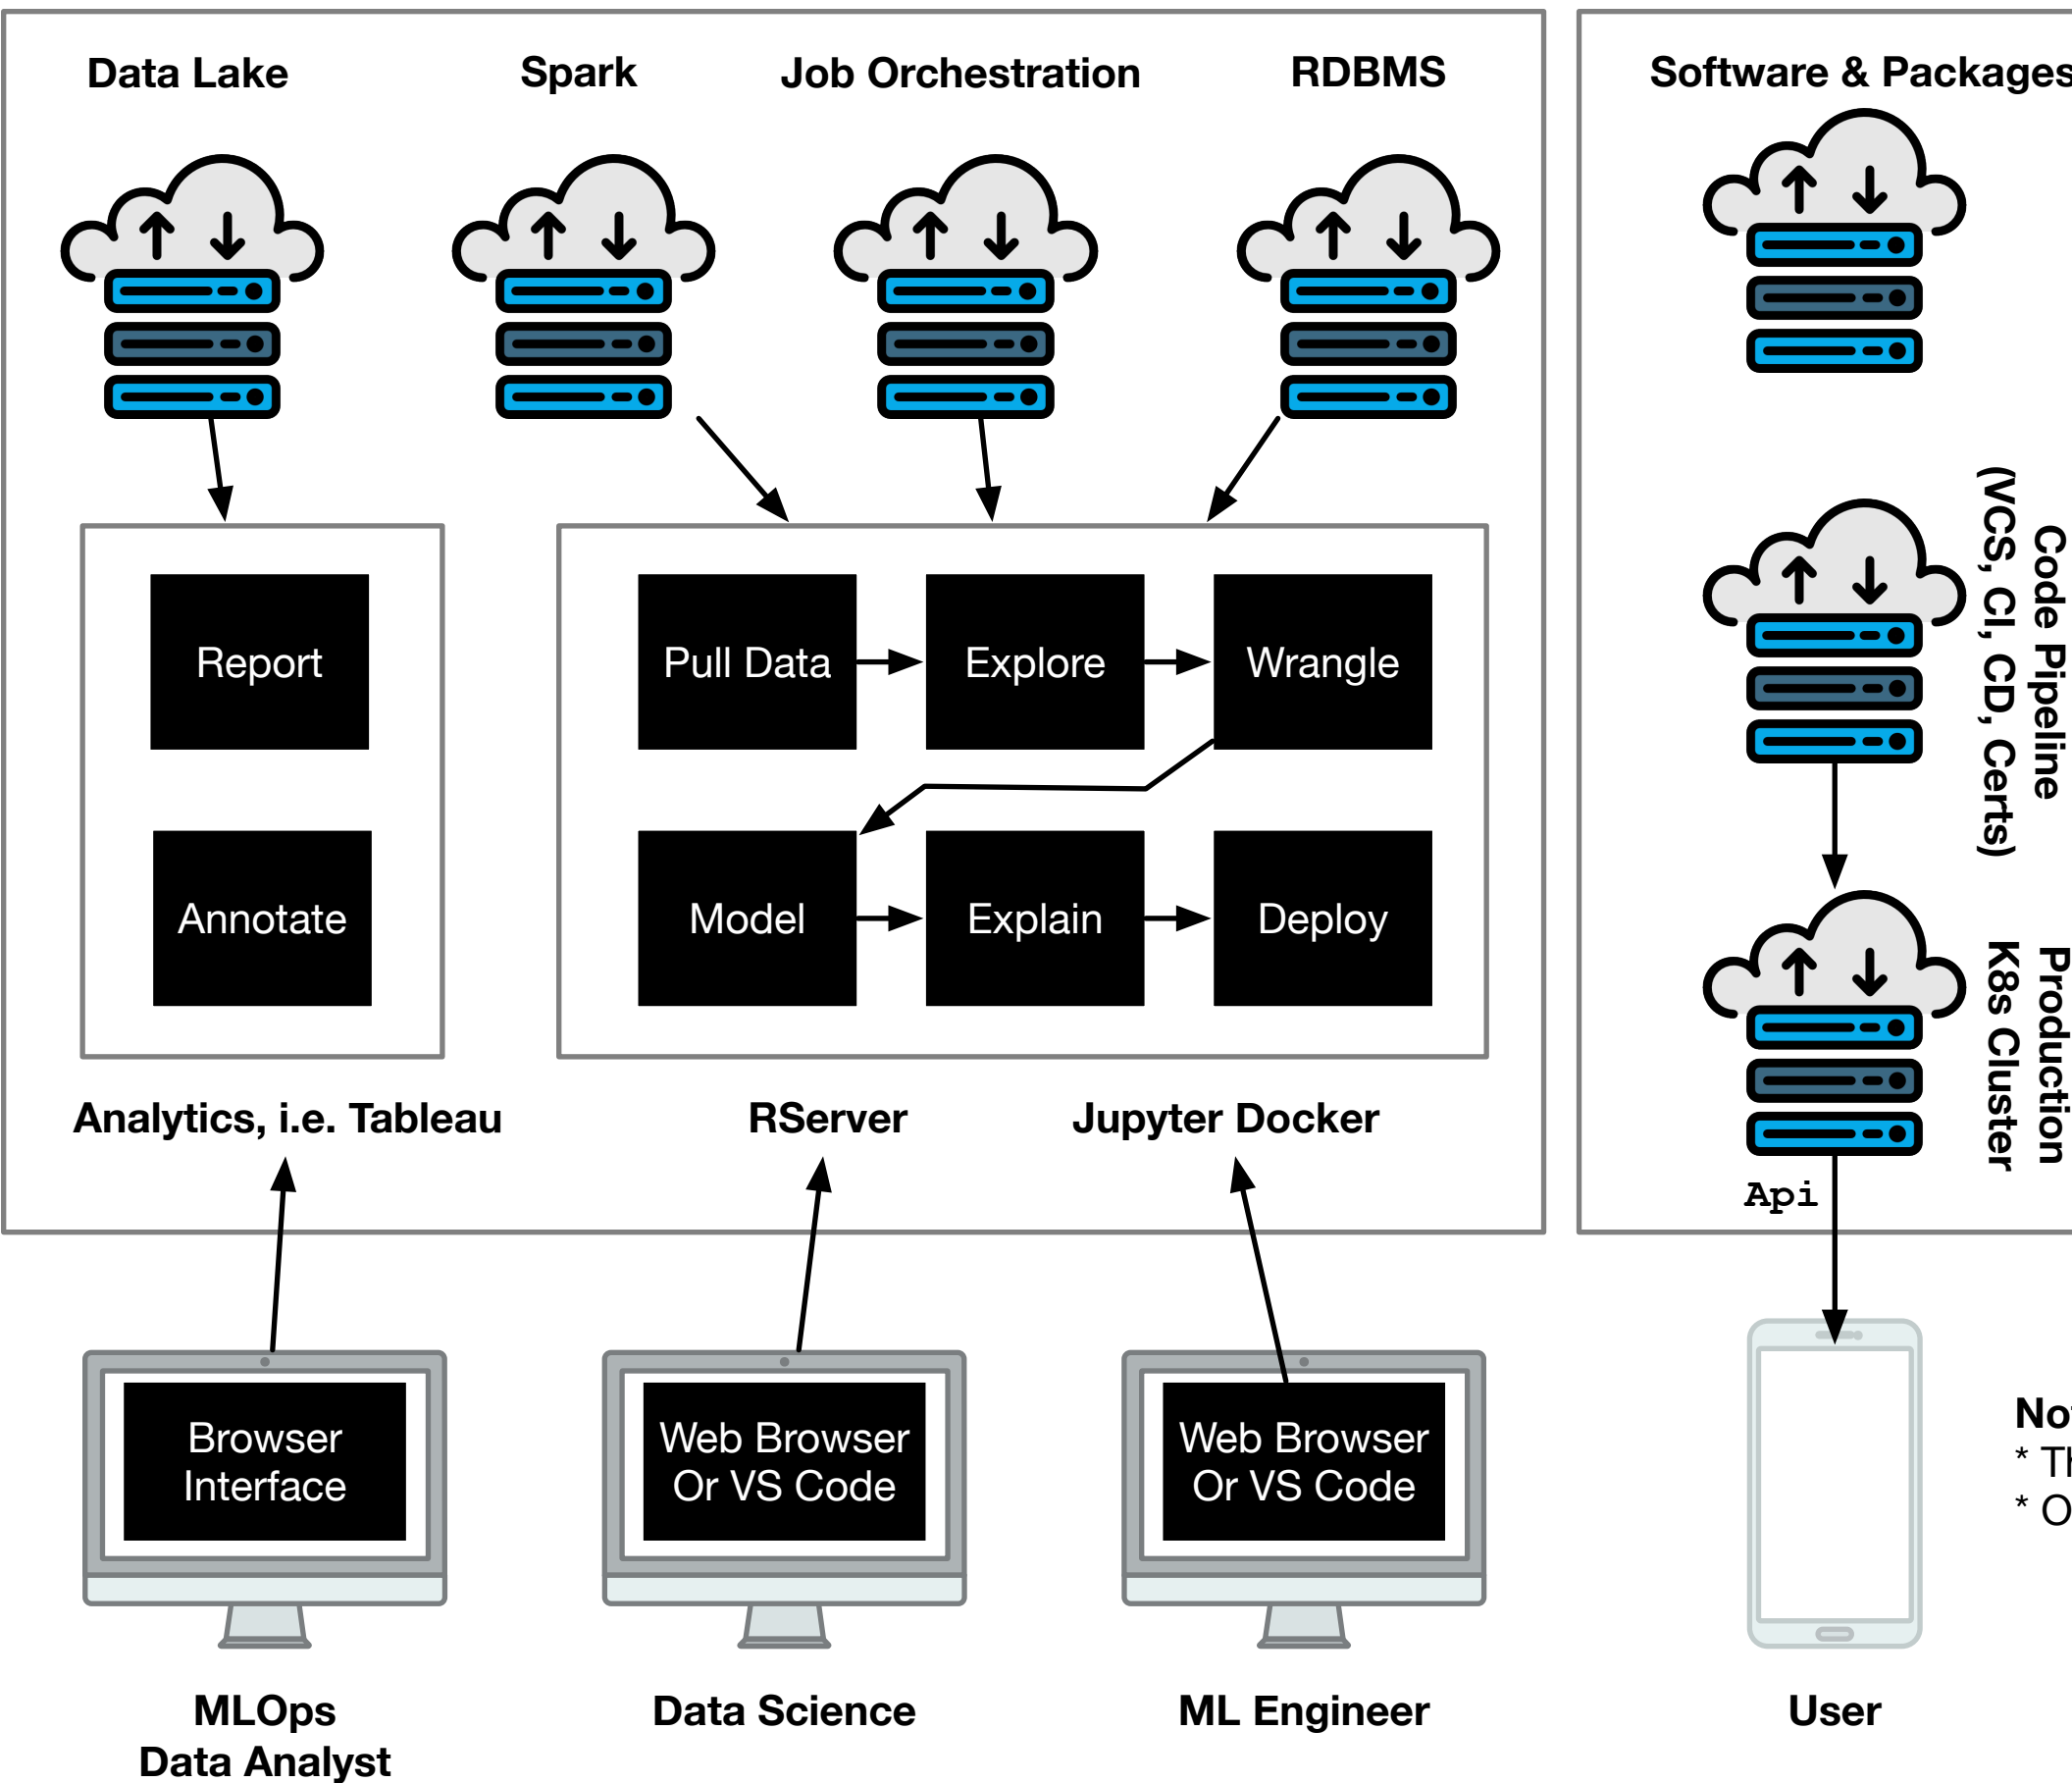
Capture Telemetry

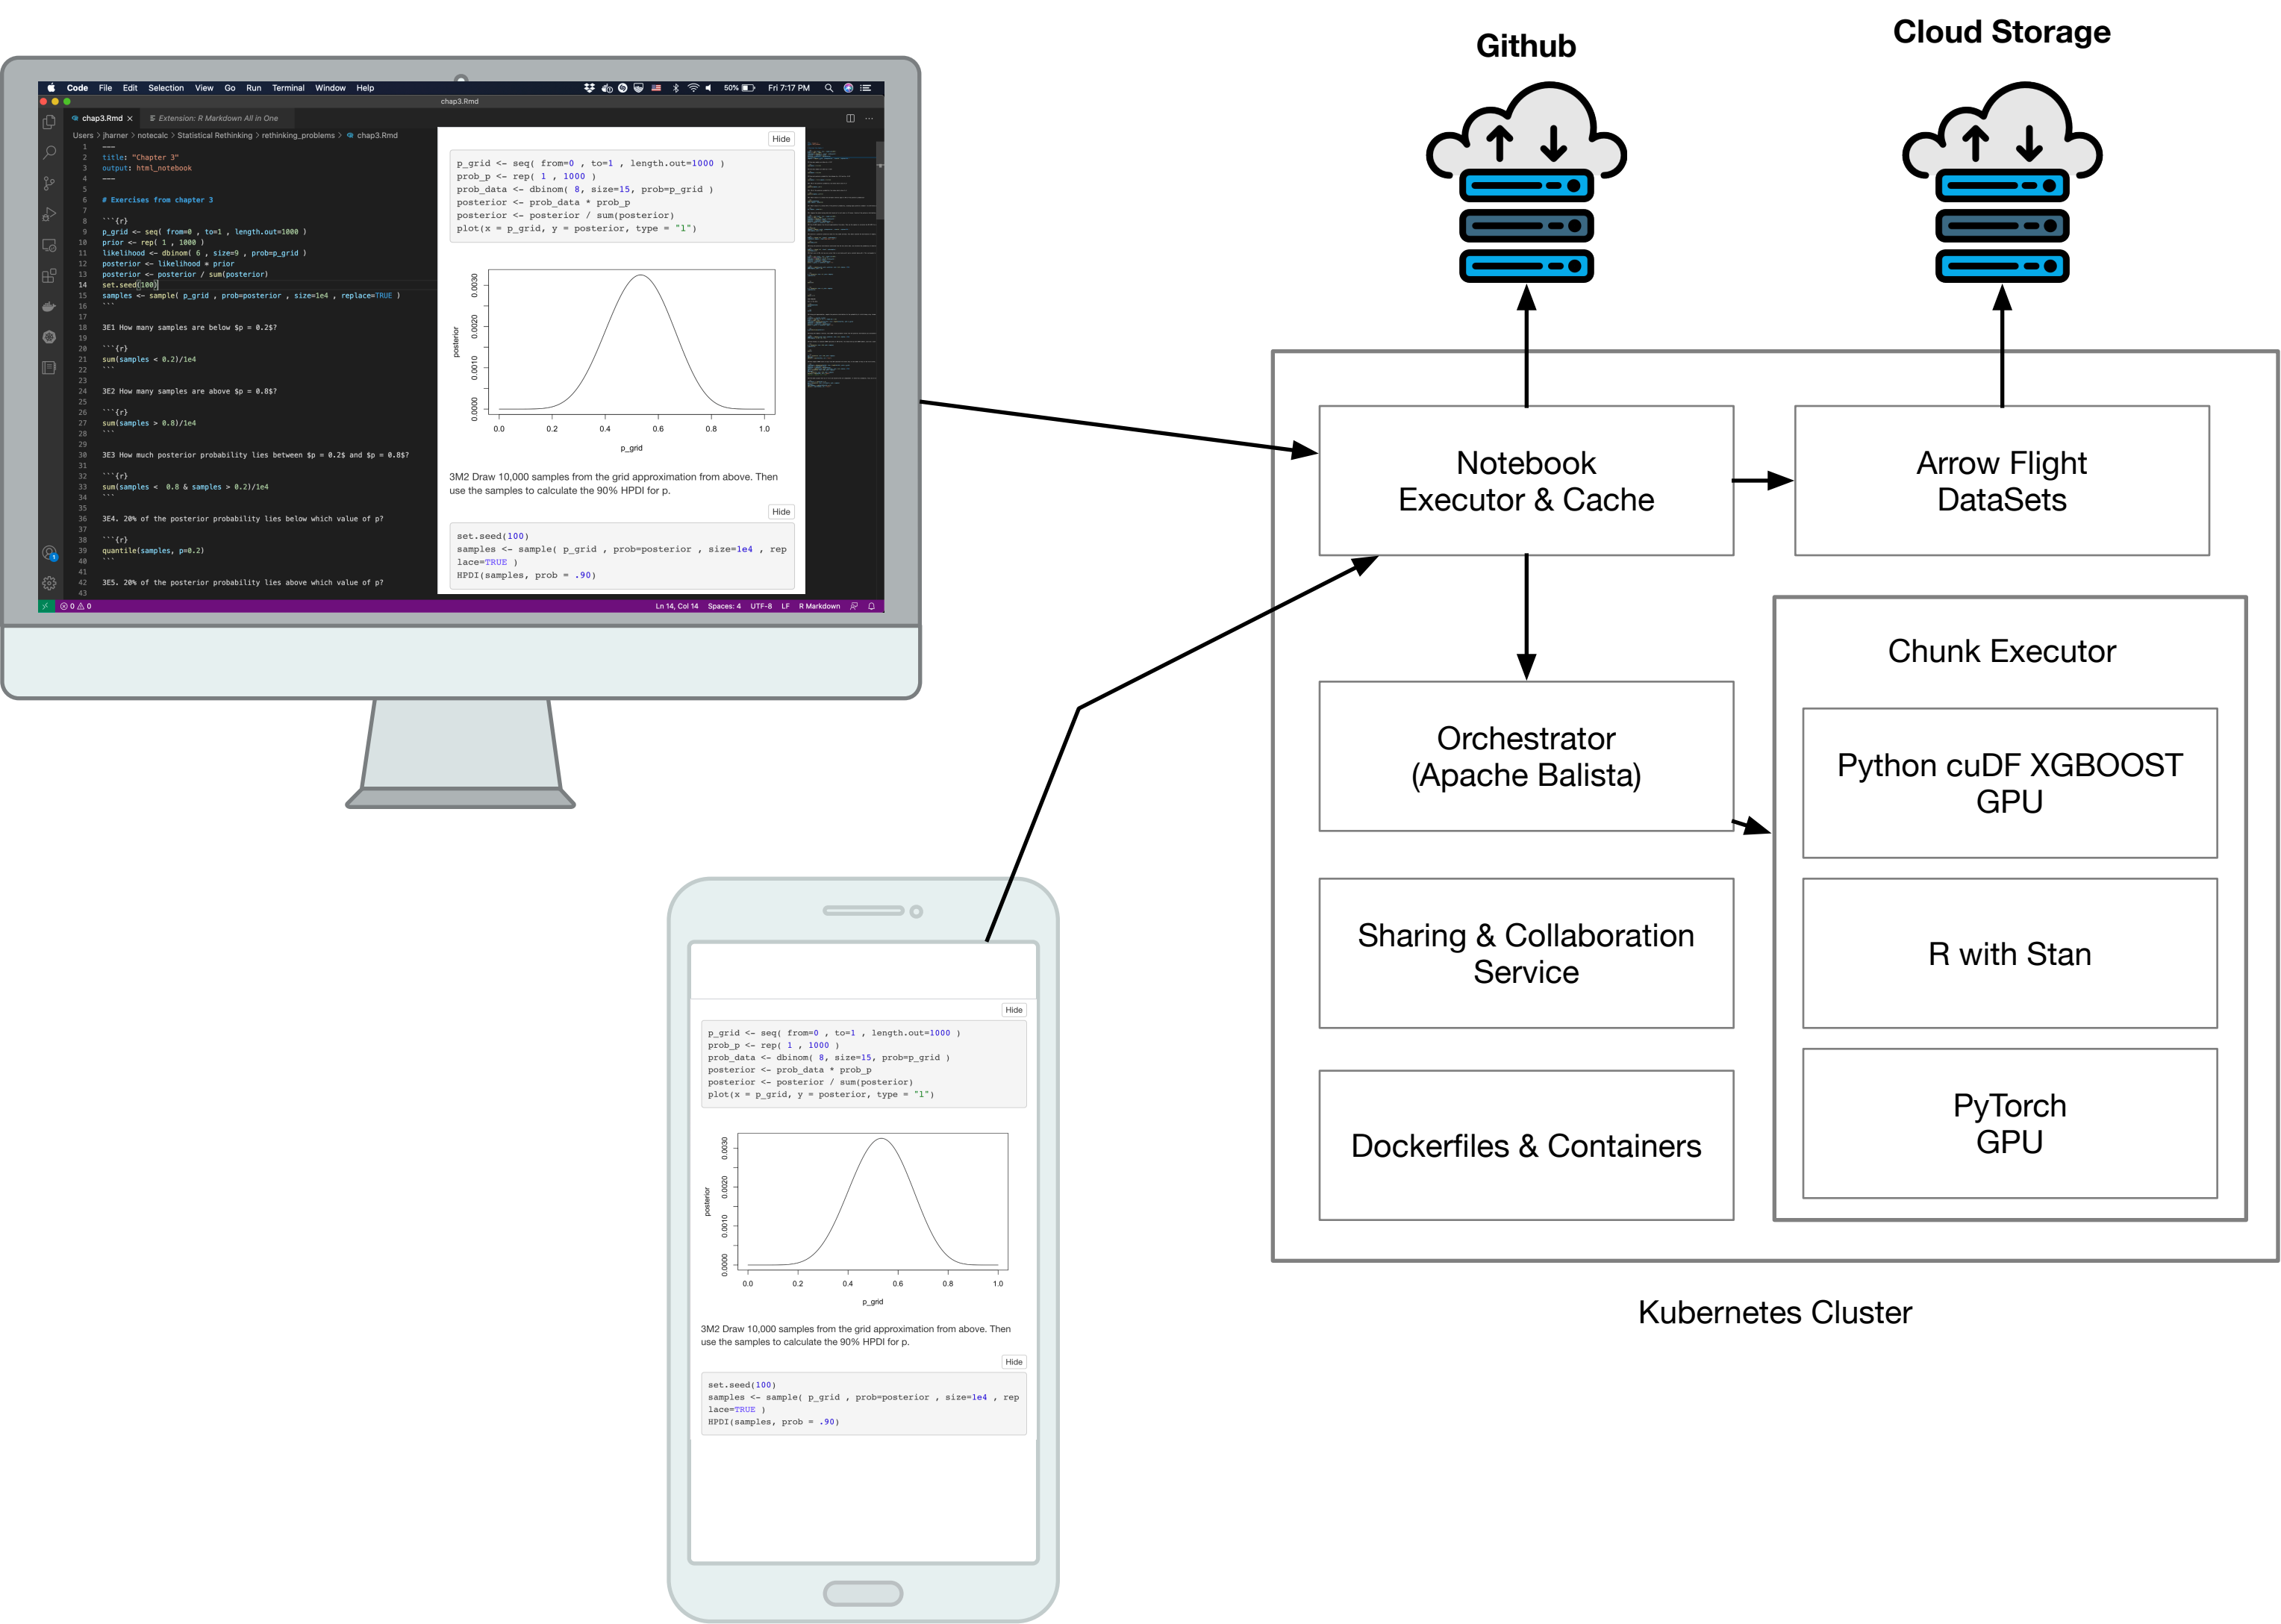
Notes:

- * How many teams work
- * Requires dev skills
- * Package hell
- * Poor Reproducibility
- * Bound to desktop tools
- * Limited compute

DataOps

DevOps








Book > Chap 3

Edit



```
prob_p <- rep( 1 , 1000 )
prob_data <- dbinom( 8, size=15, prob=p_grid )
posterior <- prob_data * prob_p
posterior <- posterior / sum(posterior)
plot(x = p_grid, y = posterior, type = "l")
```



3M2 Draw 10,000 samples from the grid approximation from above. Then use the samples to calculate the 90% HPDI for p.

Hide

```
set.seed(100)
samples <- sample( p_grid , prob=posterior , size=1e4 , replace=TRUE )
HPDI(samples, prob = .90)
```

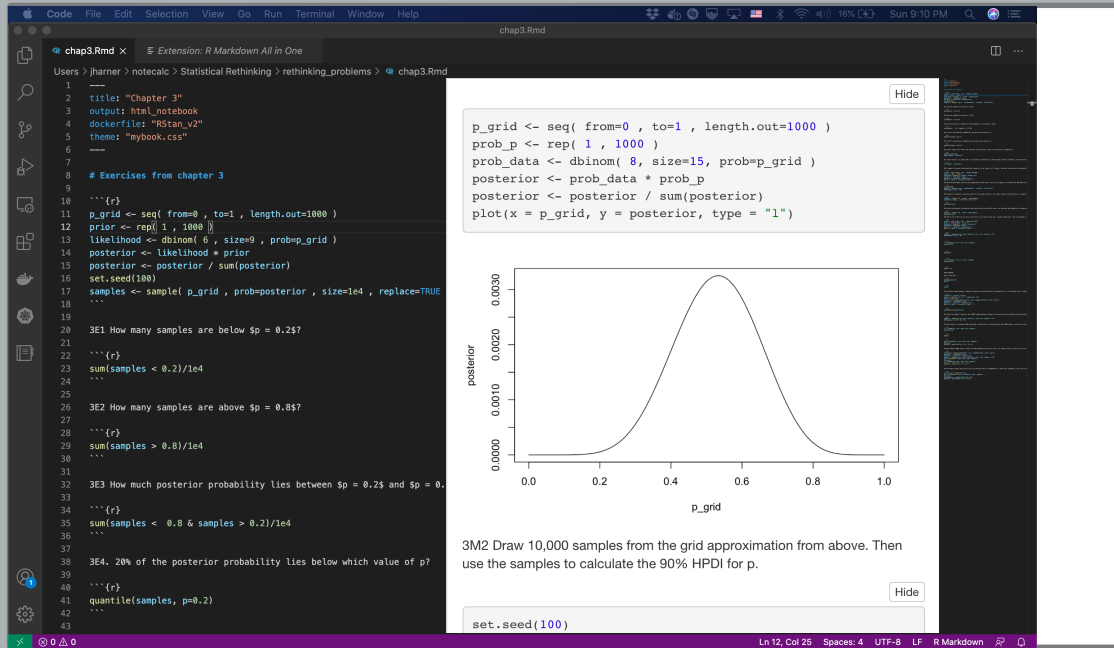
```
|0.9      0.9|
0.3343343 0.7217217
```

3M3 Construct a posterior predictive check for this model and data. This means simulate the distribution of samples, averaging over the posterior uncertainty in p. What is the probability of observing 8 water in 15 tosses?

Hide

```
dummy_w <- rbinom( 1e5 , size=15 , prob=samples )
simplehist( dummy_w , xlab="dummy water count" )
```

- * Notebook has a unique endpoint
- * Monaco Editor (from VS Code)
- * Native Application
- * Manage a set of Notebooks
- * Interactive controls for parameters
- * Edit mode
- * Share with collaborators



- * VS Code Plugin
- * Language Servers
 - syntax high-lighting
 - autocomplete
- * CSS Themes
- * Specify Docker File
- * Code blocks run on server